



Approximate Inference

Creative Machine Learning - Course 07

Pr. Philippe Esling

esling@ircam.fr



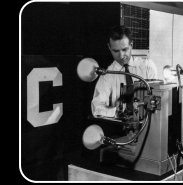
Brief history of AI

1943 - Neuron

First model by McCulloch & Pitts (purely theoretical)

1957 - Perceptron

Actual **learning machine** built by Frank Rosenblatt
Learns character recognition analogically



1986 - Backpropagation

First to learn neural networks efficiently (G. Hinton)



1989 - Convolutional NN

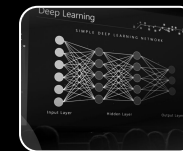
Mimicking the vision system in cats (Y. LeCun)



This lesson

2012 - Deep learning

Layerwise training to have deeper architectures
Swoop all state-of-art in classification competitions



Lesson #6

2015 - Generative model

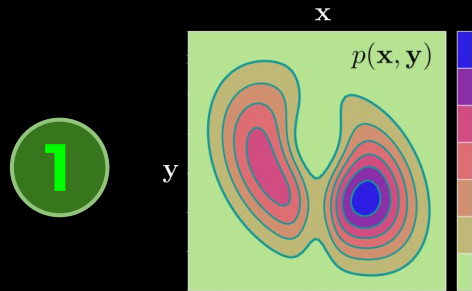
First wave of interest in generating data
Led to current model craze (VAEs, GANs, Diffusion)



2012 onwards Deep learning era

Brief history of AI

Pre-requisites for understanding generative models

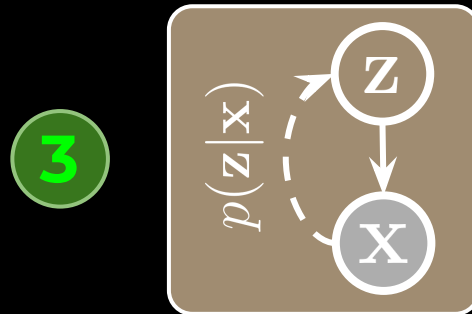
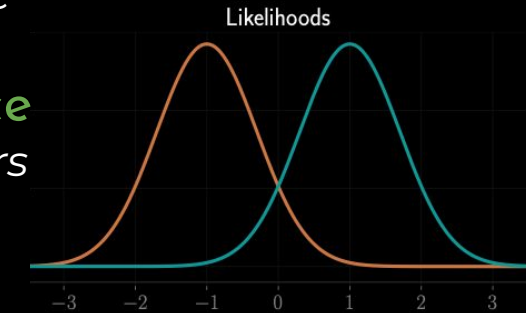


Probability theory

Random variables, distributions, independence

Bayesian inference

Bayes' theorem, likelihood, conjugate priors

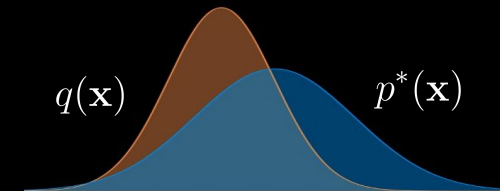


Latent models

Latent variables, probability graphs

Approximate inference

Latent variables, probability graphs



Lesson #6 2015 - Generative model

First wave of interest in generating data
Led to current model craze (VAEs, GANs, Diffusion)



2012 onwards Deep learning era

Approximate inference

Probabilistic inference deals with simple distributions

Provide easier analytical solutions

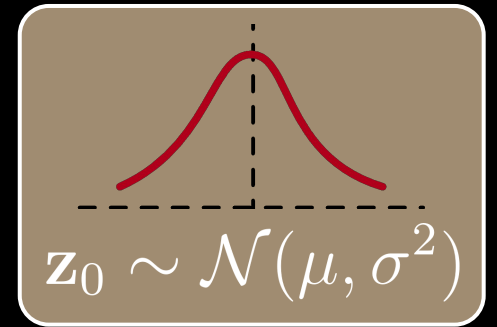
Implicit assumption of a simple explanation (capacity)

Issues

Too simplistic assumption in most cases

Real data follows largely complex distributions

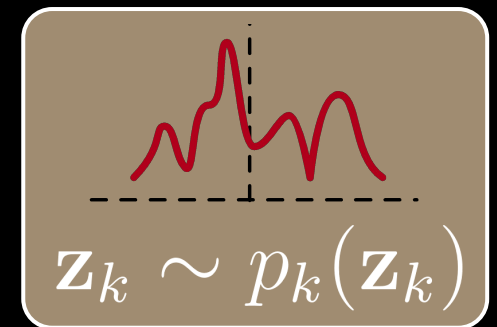
These distributions are usually intractable



The almighty Gaussian



Real data ?



How to model complex distributions but keep analytical simplicity ?

Entering the world of **approximate inference**

Approximating statistics

1. Sampling methods

Optimizing simpler distributions

1. Variational inference

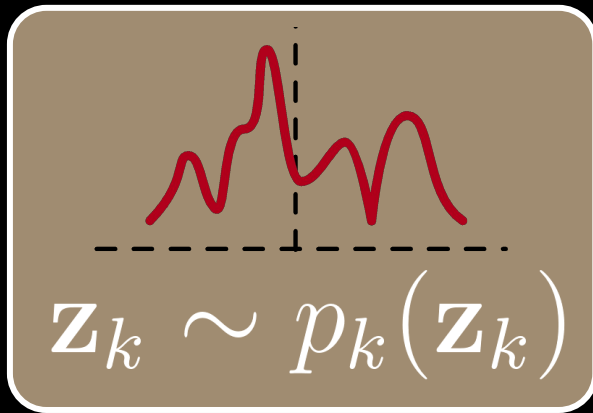
2. Normalizing flows

3. Adversarial learning

Approximate inference

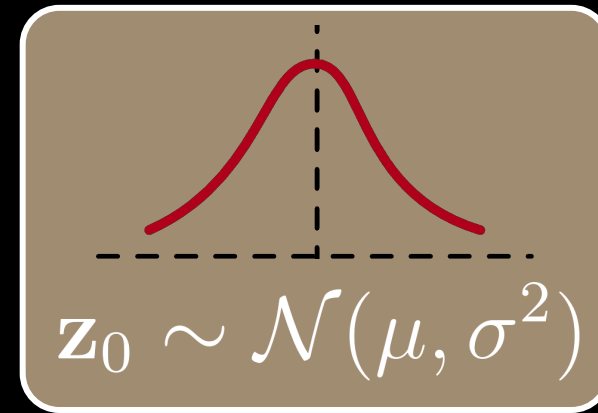
Question of the day

How can we model **real data** (of unknown distribution) with only **simple tractable distributions**



Real data ?

?



The almighty Gaussian

1st approach: can we approximate some properties ?

Examples

- Estimate the expectation $\mathbb{E}_{\mathbf{z}_k} [p_k(\mathbf{z}_k)]$
- Optimize related functions $\underset{\mathbf{z}}{\operatorname{argmax}} f(\mathbf{z}_k)$

Sampling methods

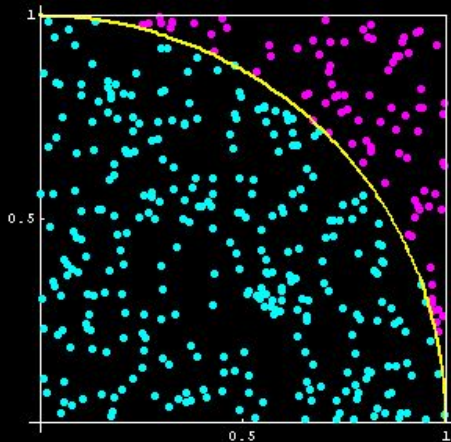
Family of *sampling methods*

- Considered the silver bullet of probabilistic modeling
- Very slow but *guaranteed to converge*
- Understand the limitations (when to use)

1 **Monte-Carlo methods**

- You have a complicated problem
- You cannot really work down the math
- Instead you can *simulate your problem* lots of time

Example Finding the value of π with a shotgun



$$\mathbb{E} [x^2 + y^2 \geq 1] \approx \frac{1}{M} \sum_{s=1}^M [x_s^2 + y_s^2 \geq 1]$$

$$\text{with } x_s, y_s \sim \mathcal{U}(0, 1)$$

Monte-Carlo methods

Goal of Monte-Carlo methods

Estimate expected values by sampling

We define an estimator $\mathbb{E}_{p(\mathbf{x})} [f(\mathbf{x})] \approx \frac{1}{M} \sum_{s=1}^M f(\mathbf{x}_s)$
with $\mathbf{x}_s \sim p(\mathbf{x})$

Interesting properties

- This is an *unbiased* estimator
- Guarantees on the convergence

Then, we can **use this estimator** in more complex problems
For instance the **M-step** of the EM algorithm

$$\max_{\theta} \mathbb{E}_q [\log p(\mathbf{x}, \mathbf{z} | \theta)]$$

Monte-Carlo methods

Advantages

- Very simple to program
- Universally applicable to lots of problems
- Scalable to parallelization

Problems

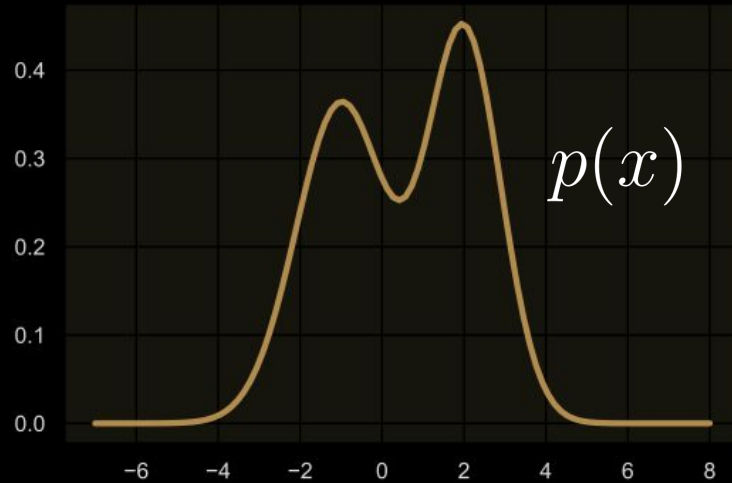
- Usually very slow (need lots of samples)
- Better solutions might exist

Major families of *sampling methods*

1. Rejection sampling
2. Markov Chain Monte-Carlo
 - Gibbs sampling
 - Metropolis-Hastings

Rejection sampling

Imagine that we have a *complex distribution* $p(x)$



We can upper-bound our distribution
Using approximation $q(x) = \mathcal{N}(1, 3)$

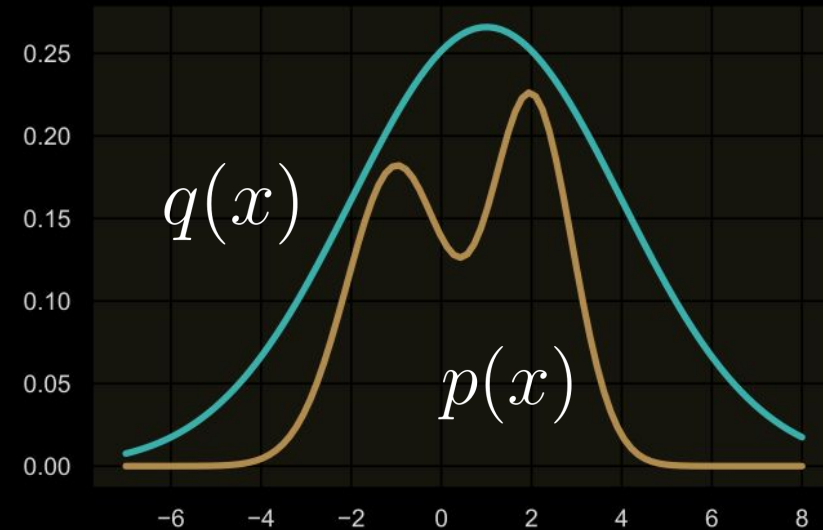
$$p(x) \leq 2q(x)$$

We use constant 2 to truly upper bound
(upper bound distribution with multiplicative constant).

We can sample from $\hat{x} \sim q(x)$

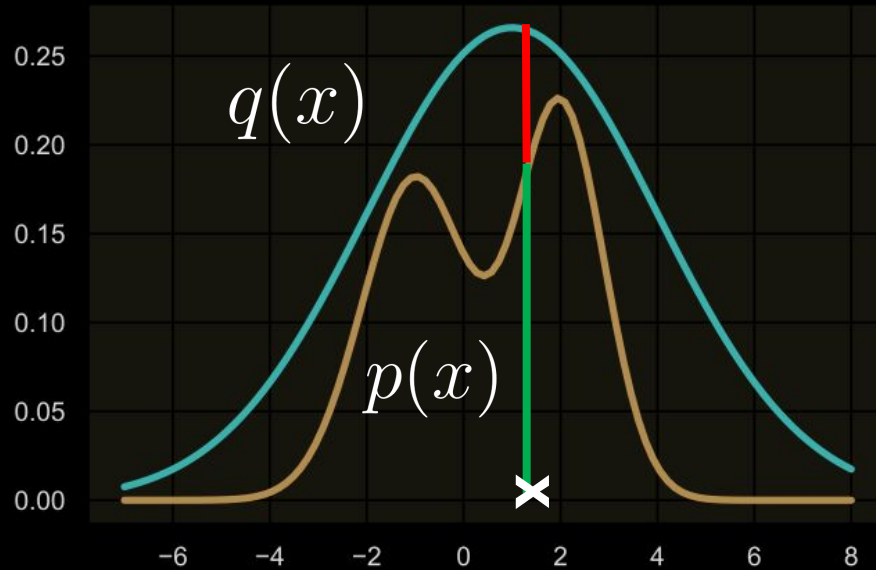
How can we sample from $p(x)$?

We know higher probability in $q(x)$
Need to ***reject*** some of the points



Rejection sampling

How do we reject some samples ?



We know it is proportional to the height

We can take a sample $\hat{x} \sim q(x)$

Take its coordinates $y \sim \mathcal{U}[0, 2q(\hat{x})]$

We accept \hat{x} with probability

$$\frac{p(x)}{2q(x)} \quad \text{if } y \geq p(x)$$

We can even use this for unnormalized distributions as

$$\frac{\hat{p}(x)}{Z} \geq Mq(x) \longrightarrow \hat{p}(x) \geq ZMq(x)$$

Markov chains

Markov chain is a sequence of random variables $\{x_0, \dots, x_n\}$
Such that x_i is independent of x_0, \dots, x_{i-2} given x_{i-1}

Markov property

$$p(x_i | x_{i-1}, \dots, x_0) = p(x_i | x_{i-1})$$

Typically described with $p(x_i | x_{i-1})$

Also noted as a *transition* $x_{k+1} \sim T(x_k \rightarrow x_{k+1})$

Stationary distribution

A distribution π is *stationnary* if it converges to a fixed distribution

$$\pi(x') = \sum_x \sim T(x \rightarrow x') \pi(x)$$

Gibbs sampling

Easiest method to *construct Markov chains to sample* from a distribution.

Base **unnormalized** distribution

$$p(x_1, x_2, x_3) = \frac{\hat{p}(x_1, x_2, x_3)}{Z}$$

Gibbs sampling

1. Start from a given **random** point (x_1^0, x_2^0, x_3^0)

2. Sample one dimension at a time

$$x_1^1 \sim p(x_1 | x_2 = x_2^0, x_3 = x_3^0) = \frac{\hat{p}(x_1, x_2^0, x_3^0)}{Z}$$

3. Re-apply this idea for next dimensions

$$x_2^1 \sim p(x_2 | x_1 = x_1^1, x_3 = x_3^0)$$

$$x_3^1 \sim p(x_3 | x_1 = x_1^1, x_2 = x_2^1)$$

Proof

$$p(x', y', z') = \sum_{x, y, z} q(x, y, z \rightarrow x', y', z') p(x, y, z)$$

Monte-Carlo methods

Advantages

- Very simple to program
- Universally applicable to lots of problems
- Scalable to parallelization

Problems

- Usually very slow (need lots of samples)
- Better solutions might exist

Sampling methods are extremely slow for complex problems

Can we turn this into an optimization problem ?

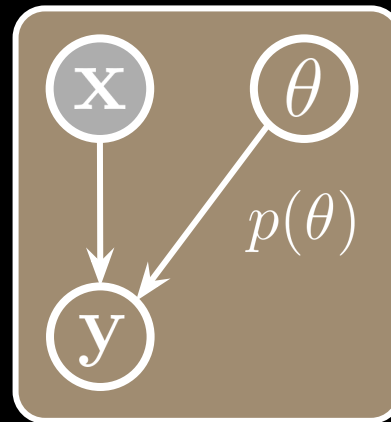
Entering the beautiful world of
variational inference

Variational inference

Goals

Variational inference approximates distributions
Optimization trick for parameter estimation
Approximation technique for *intractable integrals*
Simple example - **Bayesian logistic regression**

Data



Parameters

Full distribution

$$p(\mathbf{y}|\mathbf{x}) =$$

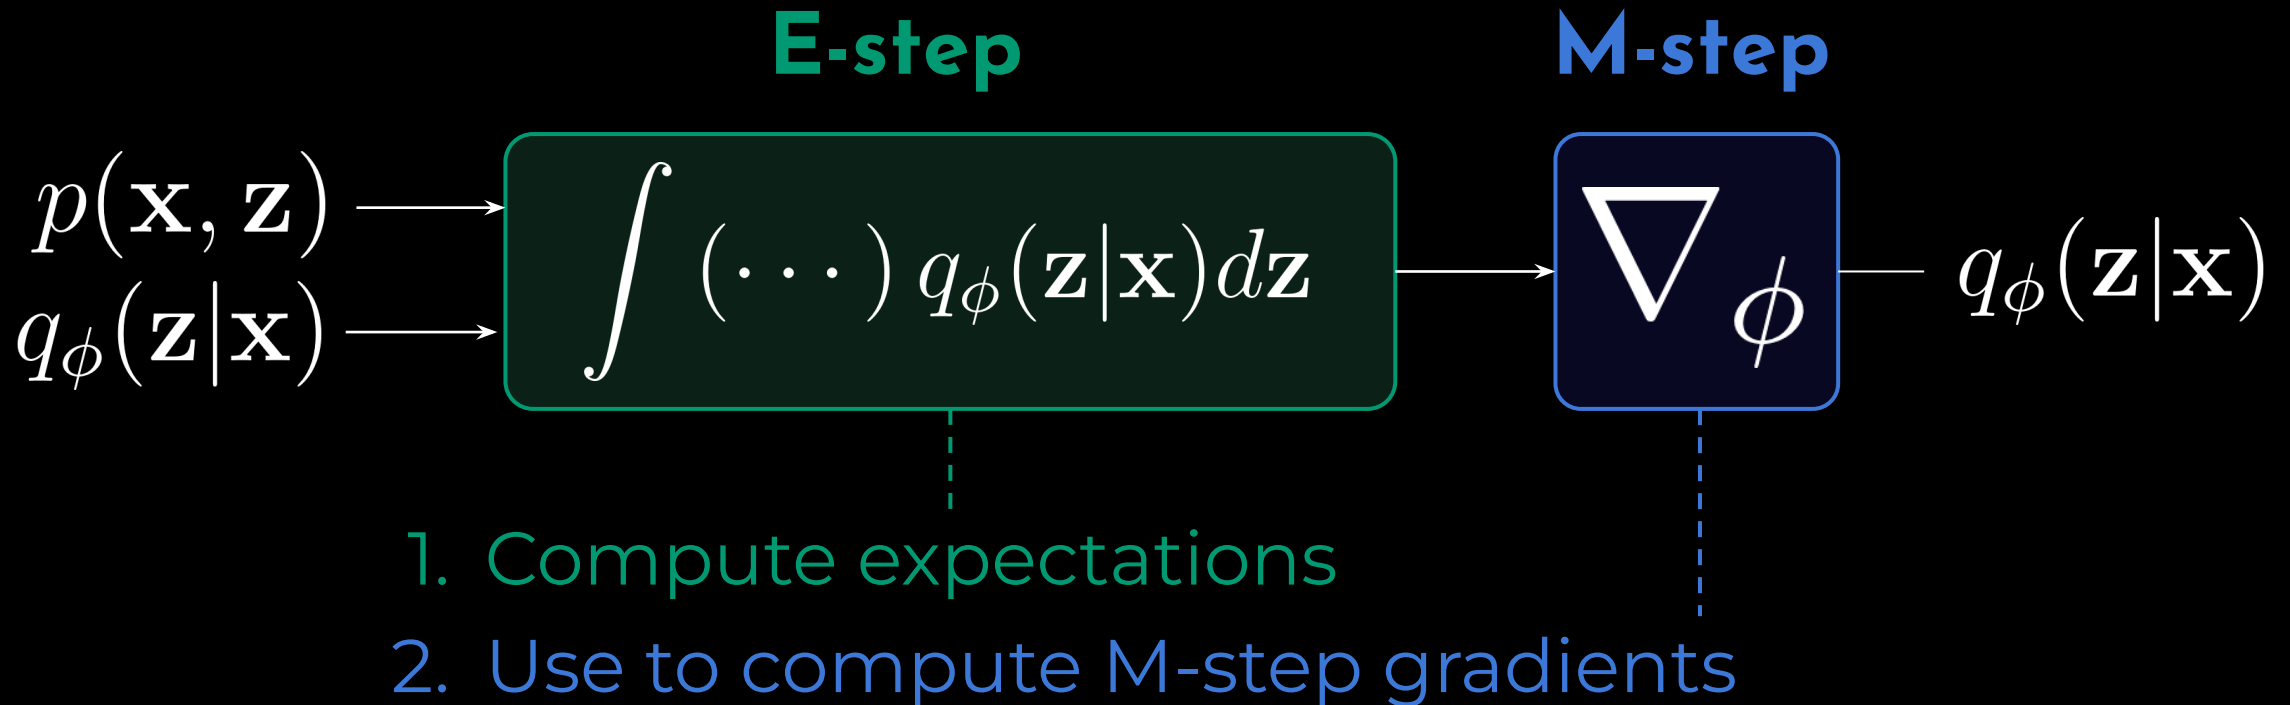
$$\int p(\mathbf{y}|\mathbf{x}, \theta) d\theta$$

Intractable integral

Classical inference approach

Recalling Expectation-Maximization

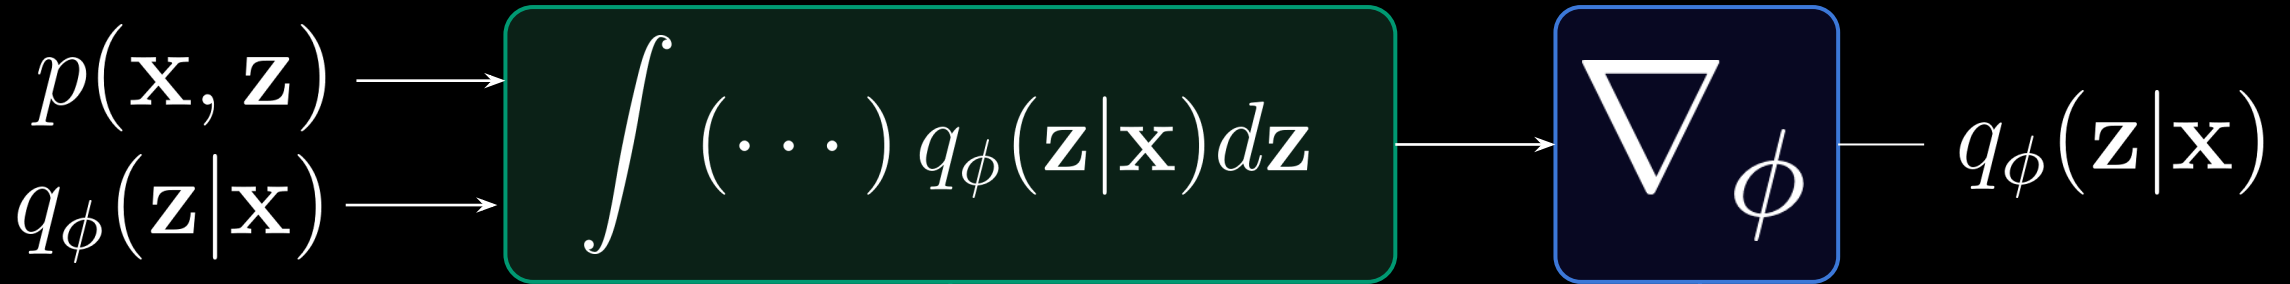
Can be summarized as the following operational flow



Stochastic inference approach

Problems

- In general, *expectations are not known*
- Gradient is of *the parameters* of the distribution
- Expectation is taken on this distribution

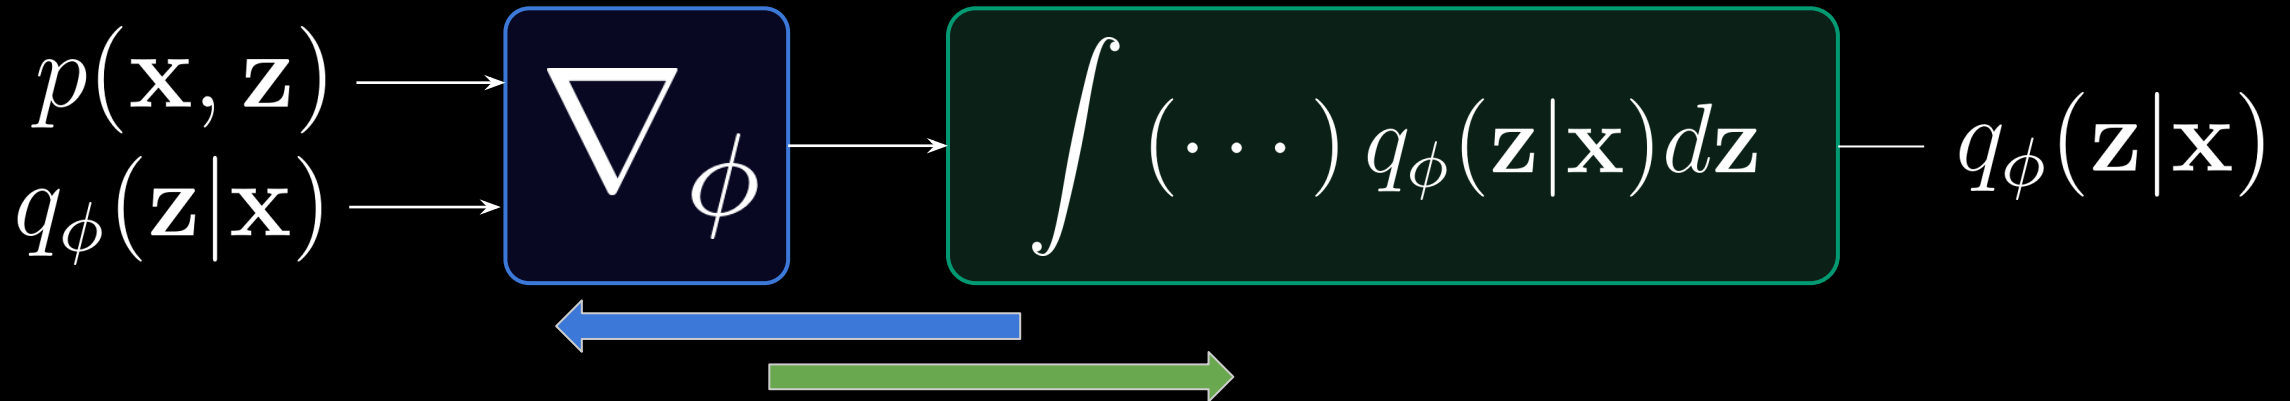


1. Compute expectations
2. Use to compute M-step gradients

Stochastic inference approach

Problems

- In general, *expectations are not known*
- Gradient is of *the parameters* of the distribution
- Expectation is taken on this distribution



Solution

- Use an *approximate* distribution
- Perform *optimization* on its parameters
- Then we can compute expectation

Stochastic gradient estimator

All approaches can be summarized as

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

Two potential solutions

Score-function estimator

Differentiate the density $q_{\phi}(\mathbf{z})$

Pathwise gradient estimator

Differentiate the function $f_{\theta}(\mathbf{z})$

Score-function estimator

If we focus on differentiating the **density**

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z}) \nabla_{\phi} \log q_{\phi}(\mathbf{z})]$$

Gradient reweighted by the value of the function

Other names

- Likelihood-ratio trick
- REINFORCE algorithm
- Automated inference
- Black-box inference

When to use

- Function is not differentiable
- Density is easy to sample from
- Density is known and differentiable

Pathwise gradient estimator

If we focus on differentiating the **function**

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

$$= \mathbb{E}_{p(\epsilon)} [\nabla_{\phi} f(g_{\phi}(\mathbf{x}, \epsilon))]$$

Introducing *reparameterization*

Find an invertible function that expresses the input as a transformation of a base distribution $\mathbf{z} = g_{\phi}(\epsilon)$ $\epsilon \sim p(\epsilon)$

$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [f_{\theta}(\mathbf{z})] = \mathbb{E}_{p(\epsilon)} [f(g_{\phi}(\mathbf{x}, \epsilon))]$$

Pathwise gradient estimator

If we focus on differentiating the **function**

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

$$= \mathbb{E}_{p(\epsilon)} [\nabla_{\phi} f(g_{\phi}(\mathbf{x}, \epsilon))]$$

Other names

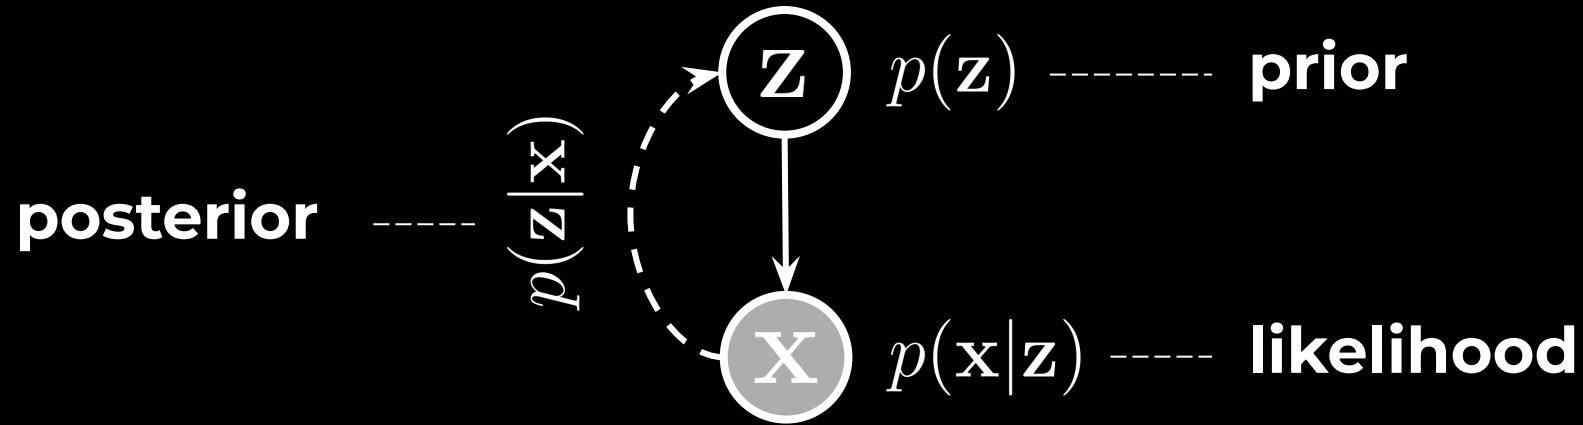
- Reparameterisation trick
- Stochastic backprop
- Perturbation analysis

When to use

- Function is differentiable
- Density can be approximated
- Easy to sample from base distrib

Variational inference

We take back our inference problem $p(\mathbf{x}, \mathbf{z}) = p(\mathbf{x}|\mathbf{z})p(\mathbf{z})$



$$p(\mathbf{x}) = \int_{\mathbb{R}} p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

How can we transform this intractable integral

Optimize a simpler and tractable distribution $q(\mathbf{z})$?

Variational inference

Starting from intractable integral $p(\mathbf{x}) = \int_{\mathbb{R}} p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$

Goal of Variational Inference

Approximate this by using **optimization** instead of **derivation**.

Select a family of *parametric* distributions $q_{\phi} \in \mathcal{Q}$

Optimize the parameters ϕ in order to solve

$$\operatorname{argmin}_{\phi} \mathcal{D}_{KL} [q_{\phi} || p]$$

Modeling the joint probability

Deriving the variational objective

$$\log p(\mathbf{x}) = \log \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

$$\log p(\mathbf{x}) = \log \int p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} q(\mathbf{z}|\mathbf{x}) d\mathbf{z} \quad \text{Introducing } q_\phi \in \mathcal{Q}$$

$$\geq \int q(\mathbf{z}|\mathbf{x}) \log \left(p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right) d\mathbf{z} \quad \text{Jensen's inequality}$$
$$f(\mathbb{E}_{q_\phi(\mathbf{z})}[\mathbf{x}]) \geq \mathbb{E}_{q_\phi(\mathbf{z})}[f(\mathbf{x})]$$

$$\geq \int q(\mathbf{z}|\mathbf{x}) \log p(\mathbf{x}|\mathbf{z}) d\mathbf{z} - \int q(\mathbf{z}|\mathbf{x}) \log \left(\frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right) d\mathbf{z}$$

$$\underbrace{\int q(\mathbf{z}|\mathbf{x}) \log p(\mathbf{x}|\mathbf{z}) d\mathbf{z}}_{\mathbb{E}_{q_\phi(\mathbf{z})}[\log p_\theta(\mathbf{x}|\mathbf{z})]} \quad \underbrace{\int q(\mathbf{z}|\mathbf{x}) \log \left(\frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right) d\mathbf{z}}_{\mathcal{D}_{KL}[q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\theta(\mathbf{z})]}$$

$$\mathcal{L}(\theta, \phi) = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z})}[\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction}} - \underbrace{\mathcal{D}_{KL}[q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\theta(\mathbf{z})]}_{\text{regularization}}$$

Variational inference

Optimizing our final objective

$$\mathcal{L}(\theta, \phi) = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z})} [\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction}} - \underbrace{\mathcal{D}_{KL} [q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\theta(\mathbf{z})]}_{\text{regularization}}$$

Optimizing reconstruction

Monte-Carlo estimator $\mathbb{E}_{q_\phi(\mathbf{z})} [\log p_\theta(\mathbf{x}|\mathbf{z})] = \frac{1}{L} \sum_{l=1}^L \log p(\mathbf{x}^{(l)}|\mathbf{z})$
... too much variance and hard to train

Reparametrization trick $\mathbf{z} = g_\phi(\mathbf{x}, \epsilon)$

deterministic part stochastic part

$$\mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [f_\theta(\mathbf{z})] = \mathbb{E}_{p(\epsilon)} [f(g_\phi(\mathbf{x}, \epsilon))] = \frac{1}{L} \sum_{l=1}^L f(g_\phi(\mathbf{x}, \epsilon^{(l)}))$$

Kullback-Leibler divergence

Optimizing regularization with KL divergence

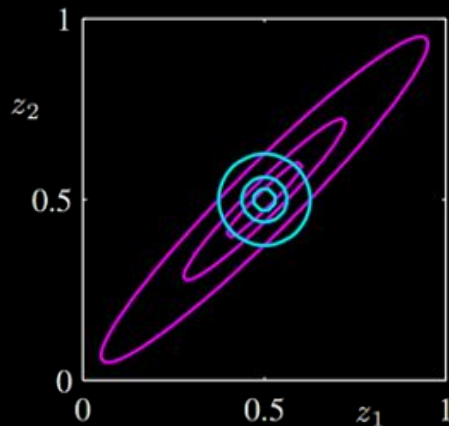
$$\mathcal{D}_{KL} [p(\mathbf{z}) \parallel q(\mathbf{z})] = \int p(\mathbf{z}) \log \frac{p(\mathbf{z})}{q(\mathbf{z})} d\mathbf{z}$$

Properties of KL

$$\mathcal{D}_{KL} [p(\mathbf{z}) \parallel q(\mathbf{z})] \geq 0. \quad \forall p, q$$

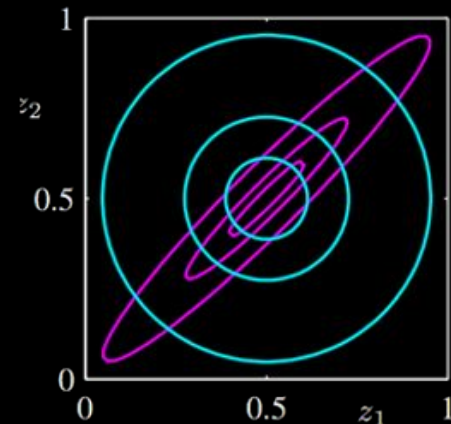
$$\mathcal{D}_{KL} [p(\mathbf{z}) \parallel q(\mathbf{z})] \neq \mathcal{D}_{KL} [q(\mathbf{z}) \parallel p(\mathbf{z})]$$

$$\mathcal{D}_{KL} [p(\mathbf{z}) \parallel q(\mathbf{z})] = 0 \text{ only if } p = q$$



Assymmetric (divergence)

$$\mathcal{D}_{KL} [p(\mathbf{z}) \parallel q(\mathbf{z})] \neq \mathcal{D}_{KL} [q(\mathbf{z}) \parallel p(\mathbf{z})]$$



Mean-field approximation

How to choose $q_\phi \in \mathcal{Q}$

flexible
easy to sample from
simple closed form

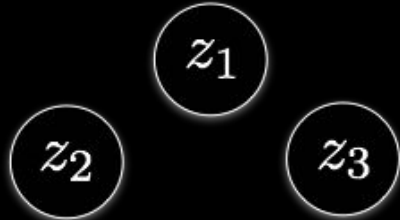
Mean-field family $q(\mathbf{z}) = \prod_{j=1}^m q_j(\mathbf{z}_j)$

- Each dimension is an independent distribution
 - Usually a Normal distribution
- Each dimension has its own factor of variation
- Compatible with conjugacy

$$p(\mathbf{z}) = p(z_1)p(z_2)p(z_3)$$

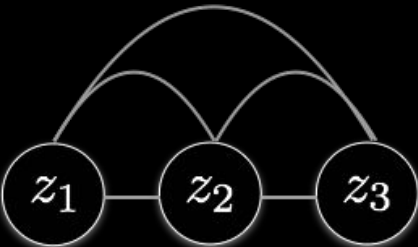
Approximation families

Least expressive



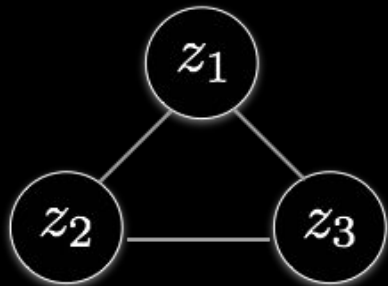
Mean-field family

$$q(\mathbf{z}) = \prod_{j=1}^m q_j(\mathbf{z}_j)$$

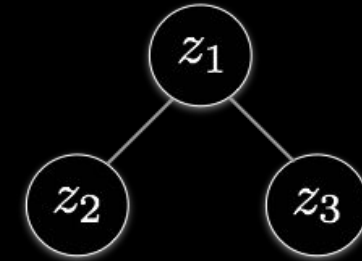


Auto-regressive

$$q(\mathbf{z}) = \prod_{i=1}^K q(z_i | z_1, \dots, z_{i-1})$$



True posterior



Structured

$$q(\mathbf{z}) = q(z_1)q(z_2|z_1)q(z_3|z_1)$$

Most expressive